

Network Analysis: Image Classification - Part 2

In the previous lab, we have implemented CNNs for an image classification task without paying much attention to the training behavior. Unchecked neural networks may or may not work based on their initialization. In this lab, we will try to exercise a fine-grain control over the parameters of the network. For this lab, you will implement the network following architecture and use the CIFAR dataset:

- conv1: convolution and rectified linear activation (RELU)
- pool1: max pooling
- conv2: convolution and rectified linear activation (RELU)
- conv3: convolution and rectified linear activation (RELU)
- pool2: max pooling
- FC1: fully connected layer with rectified linear activation (RELU)
- FC2: fully connected layer with rectified linear activation (RELU)
- FC3: fully connected layer with rectified linear activation (RELU)
- softmaxlayer: final output predictions i.e. classify into one of the ten classes.

NOTE: the kernel size, final output hidden units etc. is left up to you to design (bigger FC layers would lead to additional model complexity). You will notice that the network is a little bit poorly designed and that is intentional. We would like the network to be more complex than is needed to overfit to our data. In addition to the increased complexity of the network, we will also reduce the amount of training data to 1/2 and also remove all data augmentations/normalizations.

With this network design and data constraint, you are required to get a baseline result to showcase the model behavior. Please plot both the training and test accuracy and loss for all minibatches (not just at the end of training). This can be achieved by interleaving a testing step after every 50-100 minibatch during model training.

```
In [51]: import warnings
warnings.filterwarnings('ignore')
import torch
import torch.nn as nn
import torchvision.datasets as datasets # Standard datasets
import torchvision.transforms as transforms # Transformations we can perform on our dataset for augmentation
from torch import optim # For optimizers like SGD, Adam, etc.
from torch import nn # All neural network modules
from torch.utils.data import DataLoader, Dataset # Gives easier dataset management by creating mini batches etc.
import torch
from torch.utils.tensorboard import SummaryWriter
from torchvision import transforms
import matplotlib.pyplot as plt
torch.cuda.empty_cache()
import torch.nn.functional as F
import torchvision

import random
from timeit import default_timer as timer

In [54]: def calculateMeanStd(cifar_trainset):
    images = [item[0] for item in cifar_trainset]
    images = torch.stack(images, dim=0).numpy()

    mean_r = images[:,0,:].mean()
    mean_g = images[:,1,:].mean()
    mean_b = images[:,2,:].mean()

    # calculate std over each channel (r,g,b)
    std_r = images[:,0,:].std()
    std_g = images[:,1,:].std()
    std_b = images[:,2,:].std()

    return [mean_r, mean_g, mean_b], [std_r, std_g, std_b]

In [8]: def accuracy(outputs, labels):
    ''' Calculating accuracy of prediction'''
    _, preds = torch.max(outputs, dim=1) # class with highest probability is the predicted output
    return (torch.tensor(torch.sum(preds == labels).item() / len(preds)))

In [9]: def trainModel(model, optimizer, criterion, learning_rate, train_loader, test_loader, num_epochs, desc):
    layout = {
        "Performance on training and test data": {
            "loss": ["Multiline", ["loss/train", "loss/test"]],
            "accuracy": ["Multiline", ["accuracy/train", "accuracy/test"]],
        },
    }
    writer = SummaryWriter(f'runs/{desc}/')
    writer.add_custom_scalars(layout)

    n_total_steps = len(train_loader)
    for epoch in range(num_epochs):
        lossSum = 0
        accSum = 0

        # iterating through all batches
        for i, sample_batched in enumerate(train_loader):
            # setting gradient to zero before forward pass
            optimizer.zero_grad()

            # model predict
            yhat = model(sample_batched[0].cuda())

            # model loss
            loss = criterion(yhat.squeeze().cuda(), sample_batched[1].squeeze().cuda())
            lossSum += loss.item()

            # model accuracy
            acc = accuracy(yhat.squeeze().cuda(), sample_batched[1].squeeze().cuda())
            accSum += acc.item()

        # calculating gradient through backprop
        loss.backward(retain_graph=True)

        # updating parameters
        optimizer.step()

        # writing accuracy and loss to tensorboard logs
        lossAvg = lossSum / len(train_loader)
        accAvg = accSum / len(train_loader)
        writer.add_scalar("loss/train", lossAvg, epoch)
        writer.add_scalar("accuracy/train", accAvg, epoch)

        testAcc, testLoss = evalModel(model, optimizer, criterion, test_loader, epoch, writer)

        if (epoch+1) % 5 == 0:
            print(f'Epoch [{epoch+1}/{num_epochs}], Train Loss: {lossAvg:.4f}, Train Accuracy: {accAvg:.4f}, Test Loss: {testLoss:.4f}, Test Accuracy: {testAcc:.4f}')

        writer.flush()
        print('Finished Training')
        return model

def evalModel(model, optimizer, criterion, test_loader, epoch, writer):
    accSum = 0
    lossSum = 0
    # for evaluating model on test data with no gradient calculation
    with torch.no_grad():
        for i, sample_batched in enumerate(test_loader):
            # setting gradient to zero before forward pass
            optimizer.zero_grad()

            # model predict
            yhat = model(sample_batched[0].cuda())

            # model loss
            loss = criterion(yhat.squeeze().cuda(), sample_batched[1].squeeze().cuda())
            lossSum += loss
            # model accuracy
            acc = accuracy(yhat.squeeze().cuda(), sample_batched[1].squeeze().cuda())
            accSum += acc

        # writing accuracy and loss to tensorboard logs

    totalLoss = lossSum / len(test_loader)
    totalAccuracy = accSum / len(test_loader)
    if (writer != None):
        writer.add_scalar("loss/test", totalLoss, epoch)
        writer.add_scalar("accuracy/test", totalAccuracy, epoch)

    return totalAccuracy.item(), totalLoss.item()

In [12]: class CNN(nn.Module):
    def __init__(self, input_channels=1, num_of_classes=10):
        super(CNN, self).__init__()
        self.relu = nn.ReLU()
        self.softmax = nn.LogSoftmax()

        self.conv1 = nn.Conv2d(
            in_channels = input_channels,
            out_channels = 16,
            kernel_size = (5,5),
            stride=(1,1),
            padding=(0,0)
        )

        self.pool1 = nn.MaxPool2d(kernel_size=(2,2), stride=(2,2))

        self.conv2 = nn.Conv2d(
            in_channels = 16,
            out_channels = 32,
            kernel_size = (5,5),
            padding=(0,0)
        )

        self.conv3 = nn.Conv2d(
            in_channels = 32,
            out_channels = 64,
            kernel_size = (5,5),
            stride=(1,1),
            padding=(0,0)
        )

        self.pool2 = nn.MaxPool2d(kernel_size=(2,2), stride=(2,2))
        self.FC1 = nn.Linear(64*3*3, 1024)
        self.FC2 = nn.Linear(1024, 512)
        self.FC3 = nn.Linear(512, 256)
        self.Output = nn.Linear(256, 10)

    def forward(self, x):
        x = self.relu(self.conv1(x))
        x = self.pool1(x)
        x = self.relu(self.conv2(x))
        x = self.relu(self.conv3(x))
        x = self.pool1(x)

        x = x.view(-1, 64 * 3 * 3)

        x = self.relu(self.FC1(x))
        x = self.relu(self.FC2(x))
        x = self.relu(self.FC3(x))

        x = self.softmax(self.Output(x))

        return x

In [14]: torch.cuda.empty_cache()

# setting device
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(device)

in_channels = 3 # input channels of image
num_classes = 10 # number of output class [0-10]
num_epochs = 30
batch_size = 512

transform = transforms.Compose([transforms.Resize((32,32)),
                               transforms.ToTensor()])

# getting training and test data
train_dataset = datasets.CIFAR10(root="dataset/", train=True, transform=transform, download=True)
# splitting dataset into half with random indexes
indexList = random.sample(range(1, len(train_dataset))), (len(train_dataset)//2))
train_dataset_half = torch.utils.data.Subset(train_dataset, indexList)

test_loader = DataLoader(train_dataset_half, batch_size=batch_size, shuffle=True)

test_loader = DataLoader(train_dataset_half, batch_size=batch_size, shuffle=True)

# different setting of learning rates
learning_rate = 0.001

criterion = nn.NLLLoss()

# Initialize network
model = CNN(input_channels=in_channels, num_of_classes=num_classes).to(device)

# Initializing Optimizer
optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)

# Initializing writer for tensorboard
desc = 'Baseline'+f'_{str(num_epochs)}'

print('-----Baseline Experiment-----')
print('-----Baseline Experiment-----')

model = trainModel(model, optimizer, criterion, learning_rate, train_loader, test_loader, num_epochs, desc)

Files already downloaded and verified
In [14]: Files already downloaded and verified
In [14]: -----Baseline Experiment -----
Epoch [5/30], Train Loss: 1.5031, Train Accuracy: 0.4485, Test Loss: 1.5002, Test Accuracy: 0.4415
Epoch [10/30], Train Loss: 1.2029, Train Accuracy: 0.5685, Test Loss: 1.2735, Test Accuracy: 0.5408
Epoch [15/30], Train Loss: 0.9874, Train Accuracy: 0.6485, Test Loss: 1.2046, Test Accuracy: 0.5776
Epoch [20/30], Train Loss: 0.7771, Train Accuracy: 0.7231, Test Loss: 1.1815, Test Accuracy: 0.5988
Epoch [25/30], Train Loss: 0.5510, Train Accuracy: 0.8044, Test Loss: 1.4008, Test Accuracy: 0.5939
Epoch [30/30], Train Loss: 0.2882, Train Accuracy: 0.9018, Test Loss: 1.7907, Test Accuracy: 0.5904
Finished Training

Baseline

Performance on training and test data

loss accuracy
0 5 10 15 20 25 0 5 10 15 20 25
loss/train (Baseline_30) accuracy/train (Baseline_30)
loss/test (Baseline_30) accuracy/test (Baseline_30)

Comment: Model is overfitting training data after 15 epochs as test loss is increasing and test accuracy becomes constant and training accuracy increases.
```

Exercise 1: Normalization Effect

Now that we have a weak baseline, we can start to investigate how to improve the model performance. We will first try to address the data aspect of model improvement.

Improving data:

- Data Augmentation: the process of artificially 'increasing' our dataset by adding translation, scaling and flipping to the images to fabricate examples for training.
- Normalization: Normalizing the input data helps remove the dataset artifacts that can cause poor model performance.

In this task, you are required to add data augmentation and normalization to the dataset and run the training. Compare the training performance (accuracy and loss) to the baseline first with only data augmentation, secondly with only normalization, and lastly with both. Comment on the difference you observe in the training behavior and the final accuracy as seen on the test set. NOTE: You should use tensorboard plots.

```
In [51]: def getTrainData(augm, norm):
    # calculating mean and std for normalization
    transform = transforms.Compose([
        transforms.ToTensor()
    ])

    train_data = datasets.CIFAR10(root="dataset/",
    train=True, transform=transform, download=False)

    meanPerChannel, stdPerChannel = calculateMeanStd(train_data)

    if augm == True and norm == False:
        transform = transforms.Compose([transforms.Resize((32,32)),
        transforms.ToTensor()])

    train_dataset_org = datasets.CIFAR10(root="dataset/", train=True,
    transform=transform, download=False)

    transform = transforms.Compose([transforms.Resize((32,32)),
    transforms.ToTensor()]),
    transforms.ColorJitter(hue=0.5, saturation=0.05),
    transforms.RandomHorizontalFlip(),
    ])

    train_dataset_aug = datasets.CIFAR10(root="dataset/", train=True,
    transform=transform, download=False)

    train_dataset = torch.utils.data.ConcatDataset([train_dataset_org, train_dataset_aug])
    desc = 'with Augmentation and Normalization'

    elif augm == False and norm == True:
        transform = transforms.Compose([transforms.Resize((32,32)),
        transforms.ToTensor()]),
        transforms.Normalize(mean= meanPerChannel, std=stdPerChannel)

    train_dataset = datasets.CIFAR10(root="dataset/", train=True,
    transform=transform, download=False)

    meanPerChannel2, stdPerChannel2 = calculateMeanStd(train_dataset)
    desc = 'with Normalization Only'

    elif augm == True and norm == True:
        transform = transforms.Compose([transforms.Resize((32,32)),
        transforms.ToTensor()]),
        transforms.Normalize(mean= meanPerChannel, std=stdPerChannel)

    train_dataset_org = datasets.CIFAR10(root="dataset/", train=True,
    transform=transform, download=False)

    transform = transforms.Compose([transforms.Resize((32,32)),
    transforms.ToTensor()]),
    transforms.ColorJitter(hue=0.5, saturation=0.05),
    transforms.RandomHorizontalFlip(),
    transforms.Normalize(mean= meanPerChannel, std=stdPerChannel)

    train_dataset_aug = datasets.CIFAR10(root="dataset/", train=True,
    transform=transform, download=False)

    train_dataset = torch.utils.data.ConcatDataset([train_dataset_org, train_dataset_aug])
    desc = 'with both Augmentation and Normalization'

    return train_dataset, desc, meanPerChannel, stdPerChannel, transform

In [53]: # setting device
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

in_channels = 3 # input channels of image
num_classes = 10 # number of output class [0-10]
num_epochs = 30
batch_size = 512
config_aug_norm = [(True, False), (False, True), (True, True)]

for augm, norm in config_aug_norm:
    # getting training and test data
    train_dataset, desc, meanPerChannel, stdPerChannel, transform = getTrainData(augm, norm)
    # using same transform for test
    test_dataset = datasets.CIFAR10(root="dataset/", train=False,
    transform=transform, download=False)

    train_loader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True)
    test_loader = DataLoader(test_dataset, batch_size=batch_size, shuffle=True)

    # different setting of learning rates
    learning_rate = 0.001
    n_total_steps = len(train_loader)
    criterion = nn.NLLLoss()

    # Initialize network
    model = CNN(input_channels=in_channels, num_of_classes=num_classes).to(device)

    # Initializing Optimizer
    optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)

    # Initializing writer for tensorboard

    print('-----Experiment-----')
    print('-----Experiment-----')

    model = trainModel(model, optimizer, criterion, learning_rate, train_loader, test_loader, num_epochs, desc)

    -----Experiment with Augmentation Only-----
Epoch [5/30], Train Loss: 1.0352, Train Accuracy: 0.6310, Test Loss: 1.0736, Test Accuracy: 0.6222
Epoch [10/30], Train Loss: 0.6949, Train Accuracy: 0.7563, Test Loss: 0.9453, Test Accuracy: 0.6748
Epoch [15/30], Train Loss: 0.4947, Train Accuracy: 0.8665, Test Loss: 0.9337, Test Accuracy: 0.7151
Epoch [20/30], Train Loss: 0.2583, Train Accuracy: 0.9352, Test Loss: 1.2546, Test Accuracy: 0.6787
Epoch [25/30], Train Loss: 0.1595, Train Accuracy: 0.9483, Test Loss: 1.4912, Test Accuracy: 0.6806
Epoch [30/30], Train Loss: 0.0871, Train Accuracy: 0.9611, Test Loss: 1.7120, Test Accuracy: 0.6782
Finished Training

-----Experiment with Normalization Only-----
Epoch [5/30], Train Loss: 0.9889, Train Accuracy: 0.6494, Test Loss: 1.0373, Test Accuracy: 0.6287
Epoch [10/30], Train Loss: 0.6156, Train Accuracy: 0.7445, Test Loss: 0.9374, Test Accuracy: 0.6813
Epoch [15/30], Train Loss: 0.2989, Train Accuracy: 0.8943, Test Loss: 1.1615, Test Accuracy: 0.6884
Epoch [20/30], Train Loss: 0.1944, Train Accuracy: 0.9346, Test Loss: 1.1733, Test Accuracy: 0.7162
Epoch [25/30], Train Loss: 0.1218, Train Accuracy: 0.9609, Test Loss: 1.2723, Test Accuracy: 0.7163
Epoch [30/30], Train Loss: 0.0875, Train Accuracy: 0.9712, Test Loss: 1.4858, Test Accuracy: 0.7129
Epoch [30/30], Train Loss: 0.0822, Train Accuracy: 0.9778, Test Loss: 1.6161, Test Accuracy: 0.7214
Finished Training

With Augmentation Only

Performance on training and test data

loss accuracy
0 5 10 15 20 25 0 5 10 15 20 25
loss/train (Baseline_30) accuracy/train (Baseline_30)
loss/train (with Augmentation Only) accuracy/train (with Augmentation Only)
loss/test (Baseline_30) accuracy/test (Baseline_30)
loss/test (with Augmentation Only) accuracy/test (with Augmentation Only)

Comment: With data augmentation, more data is available for training and overfitting is reduced. also test accuracy is increased in comparison with baseline.

With Normalization Only

Performance on training and test data

loss accuracy
0 5 10 15 20 25 0 5 10 15 20 25
loss/train (Baseline_30) accuracy/train (Baseline_30)
loss/train (with Normalization Only) accuracy/train (with Normalization Only)
loss/test (Baseline_30) accuracy/test (Baseline_30)
loss/test (with Normalization Only) accuracy/test (with Normalization Only)

Comment: With normalization model is trained quickly in just around 10 epochs, in later epochs it is overfitting but overall test accuracy is better than baseline.

With both Augmentation and Normalization

Performance on training and test data

loss accuracy
0 5 10 15 20 25 0 5 10 15 20 25
loss/train (Baseline_30) accuracy/train (Baseline_30)
loss/train (with both Augmentation and Normalization) accuracy/train (with both Augmentation and Normalization)
loss/test (Baseline_30) accuracy/test (Baseline_30)
loss/test (with both Augmentation and Normalization) accuracy/test (with both Augmentation and Normalization)

Comment: When both normalization and augmentation of data is combined it gives the best result. Model is trained quickly also reduces overfitting resulting in better test accuracy than baseline.

Exercise 2: Network Regularization

Having tried to fix the dataset problems, we can have a look at the network itself. Regularization techniques are useful in learning a generalizable solution and help in avoiding overfitting. For the Deep Neural Network regularization is generally achieved by using a dropout technique, L1, L2 regularization among many other techniques. Here we will look at the dropout technique.

Dropout: Adding dropout layers to a network simply means we stochastically turn off a set number of neurons in our Fully connected layer to prevent the model from overfitting on training data.

Similar to the first task, we are interested in seeing how the addition of network regularization can improve model training behavior and overall performance. You are required to add a dropout layer and report the training and testing loss and accuracy similar to the exercise above. Comment on the impact of this modification when compared to the baseline approach.
```

```
In [27]: class CNNWithDropout(nn.Module):
    def __init__(self, input_channels=1, num_of_classes=10, dropout=0.5):
        super(CNNWithDropout, self).__init__()
        self.relu = nn.ReLU()
        self.softmax = nn.LogSoftmax()
        self.dropout = dropout

        self.conv1 = nn.Conv2d(
            in_channels = input_channels,
            out_channels = 16,
            kernel_size = (5,5),
            stride=(1,1),
            padding=(0,0)
        )

        self.pool1 = nn.MaxPool2d(kernel_size=(2,2), stride=(2,2))

        self.conv2 = nn.Conv2d(
            in_channels = 16,
            out_channels = 32,
            kernel_size = (5,5),
            stride=(1,1),
            padding=(0,0)
        )

        self.conv3 = nn.Conv2d(
            in_channels = 32,
            out_channels = 64,
            kernel_size = (5,5),
            stride=(1,1),
            padding=(0,0)
        )

        self.pool2 = nn.MaxPool2d(kernel_size=(2,2), stride=(2,2))

        self.dropout1 = nn.Dropout(self.dropout)
        self.dropout2 = nn.Dropout(self.dropout)
        self.dropout3 = nn.Dropout(self.dropout)

        self.FC1 = nn.Linear(256*3*3, 512)
        self.FC2 = nn.Linear(512, 256)
        self.FC3 = nn.Linear(256, 10)

    def forward(self, x):
        x = self.relu(self.conv1(x)) # dimension after : [1, 6, 28, 28]
        x = self.pool1(x) # dimension after : [1, 6, 14, 14]
        x = self.relu(self.conv2(x))
        x = self.relu(self.conv3(x)) # dimension after : [1, 16, 10, 10]
        x = self.pool1(x)
        # print(x.shape)
        # Flattening dimension (120, 1, 1) to 120
        x = x.view(-1, 256 * 3 * 3) # dimension after : [1, 120]

        x = self.dropout1(x)
        x = self.relu(self.FC1(x))
        x = self.dropout2(x)
        x = self.relu(self.FC2(x))
        x = self.dropout3(x)
        x = self.relu(self.FC3(x)) # dimension after : [1, 84]
        x = self.softmax(x) # dimension after : [1, 10]

        return x
```



```
# setting device
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

in_channels = 3 # input channels of image
num_classes = 10 # number of output class [0-10]
num_epochs = 30
batch_size = 32

transform = transforms.Compose([transforms.Resize((32,32)),
                               transforms.ToTensor()])

# getting training and test data
train_dataset = datasets.CIFAR10(root='dataset/', train=True, transform=transform, download=True)
test_dataset = datasets.CIFAR10(root='dataset/', train=False, transform=transform, download=True)
train_loader = DataLoader(dataset=train_dataset, batch_size=batch_size, shuffle=True)
test_loader = DataLoader(dataset=test_dataset, batch_size=batch_size, shuffle=True)

# different setting of learning rates
learning_rate=0.001
n_total_steps = len(train_loader)
criterion = nn.NLLLoss()
dropoutList = [0.3,0.4,0.5]

for dropout in dropoutList:
    desc = '%s' % str(dropout)
    # Initialize network
    model = CNNWithDropout(input_channels=in_channels, num_of_classes=num_classes, dropout =
    dropout).to(device)

    # Initializing Optimizer
    optimizer= torch.optim.Adam(model.parameters(), lr=learning_rate)

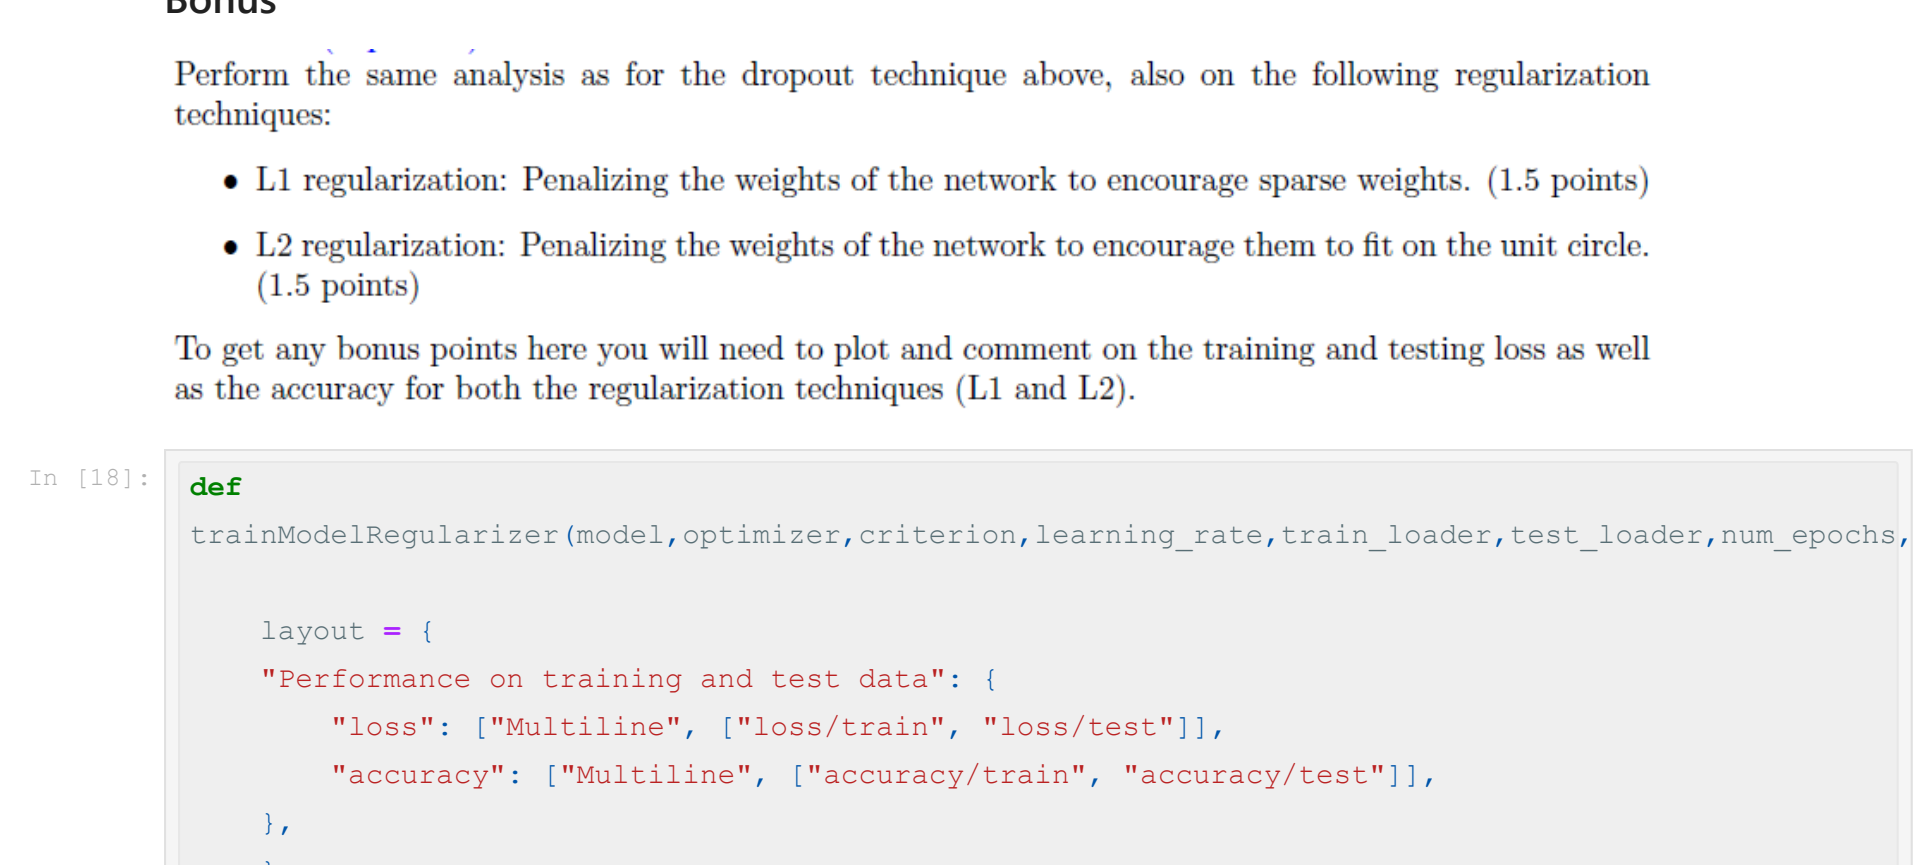
    print('----- Experiment with Dropout = %s-----' % str(dropout))
    print('-----')

    model = trainModel(model, optimizer, criterion, learning_rate, train_loader, test_loader,
    num_epochs, desc)

    writer = None
    epoch = None
    totalAccuracy = evalModel(model, optimizer, criterion, test_loader, epoch, writer)
    print('Final test Accuracy'+desc+'%s' % str(totalAccuracy))

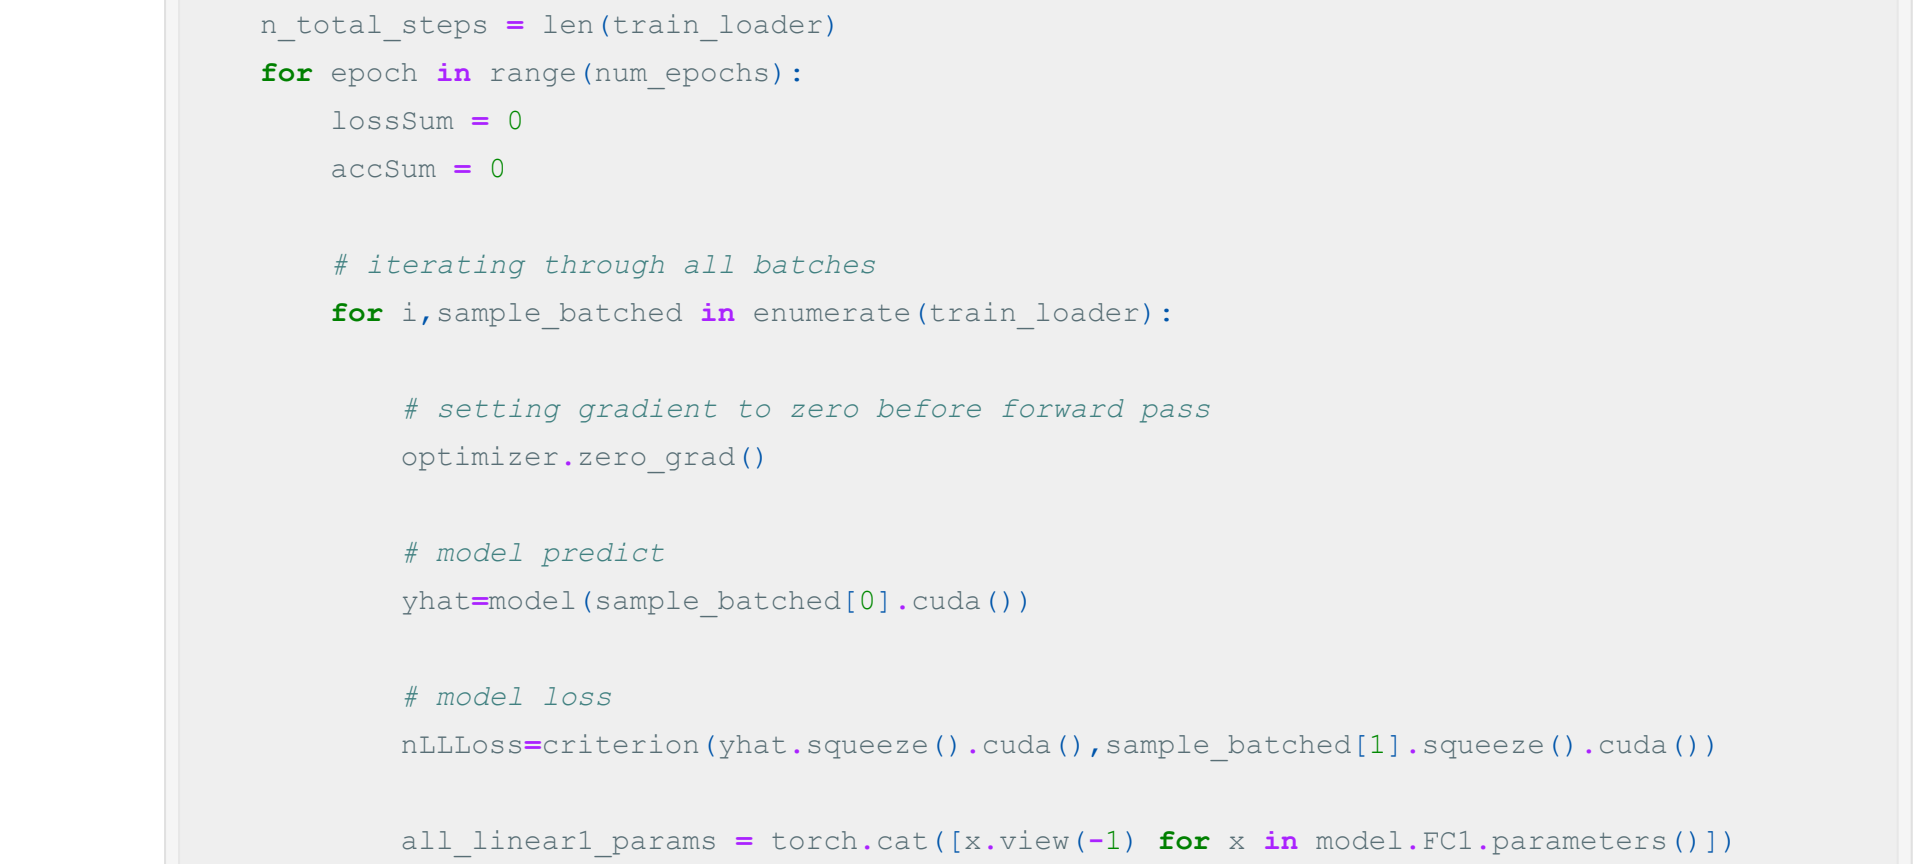
Files already downloaded and verified
Files already downloaded and verified
|----- Experiment with Dropout = 0.3-----|
Epoch [5/30], Train Loss: 1.3943, Train Accuracy: 0.4873, Test Loss: 1.3714, Test Accuracy: 0.4958
Epoch [10/30], Train Loss: 1.0138, Train Accuracy: 0.6311, Test Loss: 1.0436, Test Accuracy: 0.6252
Epoch [15/30], Train Loss: 0.7810, Train Accuracy: 0.7231, Test Loss: 0.8894, Test Accuracy: 0.6881
Epoch [20/30], Train Loss: 0.6300, Train Accuracy: 0.7808, Test Loss: 0.8289, Test Accuracy: 0.7244
Epoch [25/30], Train Loss: 0.5106, Train Accuracy: 0.8203, Test Loss: 0.8725, Test Accuracy: 0.7202
Epoch [30/30], Train Loss: 0.4161, Train Accuracy: 0.8844, Test Loss: 0.8657, Test Accuracy: 0.7331
Finished Training
Final test AccuracyDropout_0.3:(0.73331791229532, 0.8759014341504636)
|----- Experiment with Dropout = 0.4-----|
Epoch [5/30], Train Loss: 1.3949, Train Accuracy: 0.4842, Test Loss: 1.4028, Test Accuracy: 0.4873
Epoch [10/30], Train Loss: 1.0451, Train Accuracy: 0.7047, Test Loss: 1.0760, Test Accuracy: 0.6178
Epoch [15/30], Train Loss: 0.8420, Train Accuracy: 0.7827, Test Loss: 1.0809, Test Accuracy: 0.6766
Epoch [20/30], Train Loss: 0.6934, Train Accuracy: 0.7954, Test Loss: 0.8801, Test Accuracy: 0.7043
Epoch [25/30], Train Loss: 0.5851, Train Accuracy: 0.7974, Test Loss: 0.8336, Test Accuracy: 0.7231
Epoch [30/30], Train Loss: 0.4862, Train Accuracy: 0.8287, Test Loss: 0.8815, Test Accuracy: 0.7166
Finished Training
Final test AccuracyDropout_0.4:(0.718422532081604, 0.8820583058372388)
|----- Experiment with Dropout = 0.5-----|
Epoch [5/30], Train Loss: 1.3949, Train Accuracy: 0.4892, Test Loss: 1.3792, Test Accuracy: 0.4988
Epoch [10/30], Train Loss: 1.0397, Train Accuracy: 0.6301, Test Loss: 1.0771, Test Accuracy: 0.6195
Epoch [15/30], Train Loss: 0.8339, Train Accuracy: 0.7107, Test Loss: 0.9489, Test Accuracy: 0.6761
Epoch [20/30], Train Loss: 0.6885, Train Accuracy: 0.7616, Test Loss: 0.8671, Test Accuracy: 0.7073
Epoch [25/30], Train Loss: 0.5851, Train Accuracy: 0.7974, Test Loss: 0.8336, Test Accuracy: 0.7231
Epoch [30/30], Train Loss: 0.5097, Train Accuracy: 0.8277, Test Loss: 0.8815, Test Accuracy: 0.7324
Finished Training
Final test AccuracyDropout_0.5:(0.7274242043495178, 0.8679054379461396)
```

Comparison with Baseline



Comment: When comparing with baseline, test accuracy increases after adding dropout because it regularizes the neurons in layers and model does not overfit the training data, which was the case in baseline after 15 epochs as test loss is increasing.

Comparison with different dropouts



Comment: When experimenting with different dropout rates, test accuracy is best for dropout=0.5.

Bonus

Perform the same analysis as for the dropout technique above, also on the following regularization techniques:

- L1 regularization: Penalizing the weights of the network to encourage sparse weights. (1.5 points)
- L2 regularization: Penalizing the weights of the network to encourage them to fit on the unit circle. (1.5 points)

To get any bonus points here you will need to plot and comment on the training and testing loss as well as the accuracy for both the regularization techniques (L1 and L2).

```
In (18): def trainModelRegularizer(model, optimizer, criterion, learning_rate, train_loader, test_loader, num_epochs,
                             layout = {
                                 "Performance on training and test data": {
                                     "loss": ["Multiline", ["loss/train", "loss/test"]],
                                     "accuracy": ["Multiline", ["accuracy/train", "accuracy/test"]],
                                 },
                             )
    writer = SummaryWriter('%sruns/%s' % (run_dir, desc))
    writer.add_custom_scalars(layout)

    n_total_steps = len(train_loader)
    for epoch in range(num_epochs):
        lossSum = 0
        accSum = 0

        # iterating through all batches
        for isample_batched in enumerate(train_loader):
            # setting gradient to zero before forward pass
            optimizer.zero_grad()

            # model predict
            yhat=model(torch.tensor(sample_batched[0].cuda()))

            # model loss
            nLLoss=criterion(yhat.squeeze().cuda(), sample_batched[1].squeeze().cuda())

            all_linear_params = torch.cat([x.view(-1) for x in model.fc1.parameters()])
            l1_normalization = lambda() for torch.nn.Parameter in all_linear_params, regularizer)
            #l1_norm = sum(p.abs()).sum() for l1_norm in model.parameters())

            loss = nLLoss + l1_regularization
            lossSum += loss.item()

            model accuracy
            acc = accuracy(yhat.squeeze().cuda(), sample_batched[1].squeeze().cuda())
            accSum += acc.item()

            # calculating gradient through backprop
            loss.backward(retain_graph=True)

            # updating parameters
            optimizer.step()

            # writing accuracy and loss to tensorboard logs
            lossAvg = lossSum / len(train_loader)
            accAvg = accSum / len(train_loader)
            writer.add_scalar("loss/train", lossAvg, epoch)
            writer.add_scalar("accuracy/train", accAvg, epoch)

            testAcc, testLoss = evalModel(model, optimizer, criterion, test_loader, epoch, writer)

            if (epoch+1) % 5 == 0:
                print ("Epoch [%d/%d] [%d/%d] Train Loss: {lossAvg:.4f}, Train Accuracy: {accAvg:.4f}, Test Loss: {testLoss:.4f}, Test Accuracy: {testAcc:.4f}")

            writer.flush()
            print("Finished Training")
    return model
```

```
In (19): # setting device
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(device)

in_channels = 3 # input channels of image
num_classes = 10 # number of output class [0-10]
num_epochs = 30
batch_size = 32

transform = transforms.Compose([transforms.Resize((32,32)),
                               transforms.ToTensor()])

# getting training and test data
train_dataset = datasets.CIFAR10(root='dataset/', train=True, transform=transform, download=True)
# splitting dataset into half with random indexes
indexList = random.sample(list(range(0,len(train_dataset))), (len(train_dataset)//2))
train_dataset_half = torch.utils.data.Subset(train_dataset, indexList)

test_dataset = datasets.CIFAR10(root='dataset/', train=False, transform=transform, download=True)
train_loader = DataLoader(dataset=train_dataset_half, batch_size=batch_size, shuffle=True)
test_loader = DataLoader(dataset=test_dataset, batch_size=batch_size, shuffle=True)

# different setting of learning rates
learning_rate=0.001
n_total_steps = len(train_loader)
criterion = nn.NLLLoss()
for regularizer, lambda in config:
    # Initialize network
    model = CNN(input_channels=in_channels, num_of_classes=num_classes).to(device)

    # Initializing Optimizer
    optimizer= torch.optim.Adam(model.parameters(), lr=learning_rate)

    # Initializing writer for tensorboard
    desc = '%s' % str(regularizer) + '%s' % str(lambda)

    print('----- Experiment with %s-----' % str(lambda))
    print('-----')

    model = trainModelRegularizer(model, optimizer, criterion, learning_rate,
    train_loader, test_loader, num_epochs, desc, regularizer, lambda)

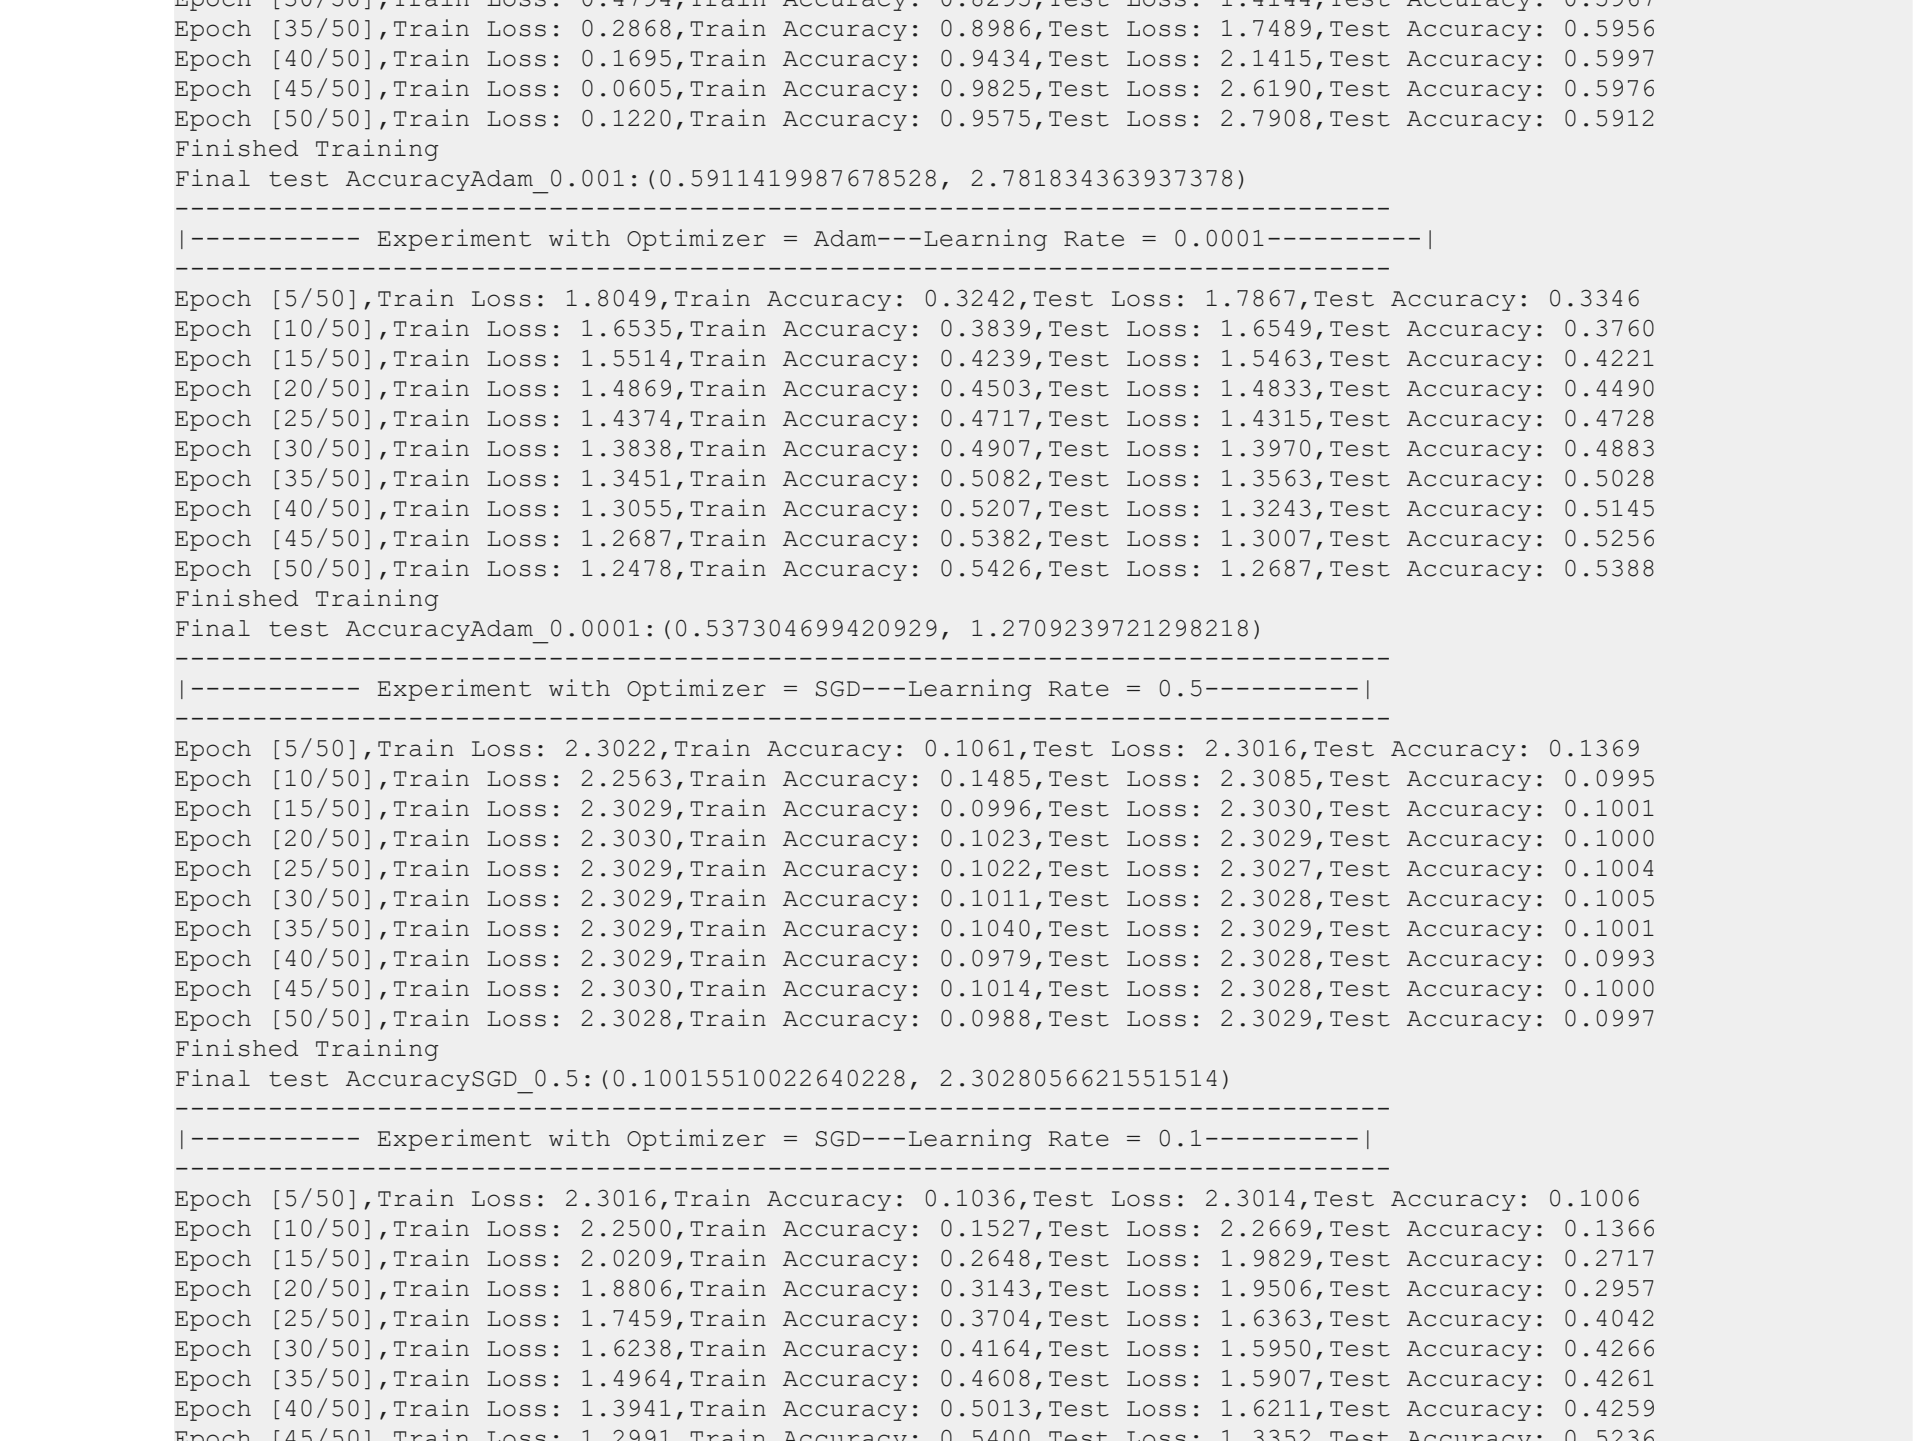
cuda
Files already downloaded and verified
Files already downloaded and verified
|-----L1_0.05 Experiment-----|
Epoch [5/30], Train Loss: 1.3933, Train Accuracy: 0.1039, Test Loss: 1.3922, Test Accuracy: 0.1039
Epoch [10/30], Train Loss: 1.2505, Train Accuracy: 0.2108, Test Loss: 1.3218, Test Accuracy: 0.2018
Epoch [15/30], Train Loss: 1.0929, Train Accuracy: 0.2288, Test Loss: 1.1809, Test Accuracy: 0.2308
Epoch [20/30], Train Loss: 0.9309, Train Accuracy: 0.2445, Test Loss: 1.1531, Test Accuracy: 0.2376
Epoch [25/30], Train Loss: 0.8321, Train Accuracy: 0.2456, Test Loss: 1.1531, Test Accuracy: 0.2376
Epoch [30/30], Train Loss: 0.7167, Train Accuracy: 0.2540, Test Loss: 1.1470, Test Accuracy: 0.2390
Finished Training
Final test AccuracyL1_0.05:(0.254025402540254, 0.239023902390239)
|-----L1_0.01 Experiment-----|
Epoch [5/30], Train Loss: 1.6926, Train Accuracy: 0.2480, Test Loss: 1.8766, Test Accuracy: 0.2682
Epoch [10/30], Train Loss: 1.3931, Train Accuracy: 0.3474, Test Loss: 1.6723, Test Accuracy: 0.3604
Epoch [15/30], Train Loss: 1.0708, Train Accuracy: 0.4231, Test Loss: 1.5143, Test Accuracy: 0.4120
Epoch [20/30], Train Loss: 0.7224, Train Accuracy: 0.4448, Test Loss: 1.5621, Test Accuracy: 0.4038
Epoch [25/30], Train Loss: 0.5228, Train Accuracy: 0.4773, Test Loss: 1.4895, Test Accuracy: 0.4507
Epoch [30/30], Train Loss: 0.2666, Train Accuracy: 0.5038, Test Loss: 1.4707, Test Accuracy: 0.4993
Finished Training
Final test AccuracyL1_0.01:(0.50382667896747589, 0.49932667896747589)
|-----L1_0.001 Experiment-----|
Epoch [5/30], Train Loss: 1.7221, Train Accuracy: 0.3854, Test Loss: 1.6200, Test Accuracy: 0.3896
Epoch [10/30], Train Loss: 1.4895, Train Accuracy: 0.4940, Test Loss: 1.3195, Test Accuracy: 0.5008
Epoch [15/30], Train Loss: 1.3231, Train Accuracy: 0.5656, Test Loss: 1.2347, Test Accuracy: 0.5475
Epoch [20/30], Train Loss: 1.2139, Train Accuracy: 0.6150, Test Loss: 1.2196, Test Accuracy: 0.5867
Epoch [25/30], Train Loss: 1.1240, Train Accuracy: 0.6377, Test Loss: 1.1470, Test Accuracy: 0.5950
Epoch [30/30], Train Loss: 1.0341, Train Accuracy: 0.6965, Test Loss: 1.1822, Test Accuracy: 0.5844
Finished Training
Final test AccuracyL1_0.001:(0.69652667896747589, 0.58442667896747589)
|-----L2_0.01 Experiment-----|
Epoch [5/30], Train Loss: 1.5906, Train Accuracy: 0.4224, Test Loss: 1.4859, Test Accuracy: 0.4303
Epoch [10/30], Train Loss: 1.3051, Train Accuracy: 0.5449, Test Loss: 1.2393, Test Accuracy: 0.5367
Epoch [15/30], Train Loss: 1.0708, Train Accuracy: 0.6057, Test Loss: 1.2143, Test Accuracy: 0.5676
Epoch [20/30], Train Loss: 0.7999, Train Accuracy: 0.6257, Test Loss: 1.1595, Test Accuracy: 0.6002
Epoch [25/30], Train Loss: 0.5969, Train Accuracy: 0.6373, Test Loss: 1.2196, Test Accuracy: 0.5867
Epoch [30/30], Train Loss: 0.3967, Train Accuracy: 0.6732, Test Loss: 1.1745, Test Accuracy: 0.6074
Finished Training
Final test AccuracyL2_0.01:(0.67322667896747589, 0.60742667896747589)
|-----L2_0.001 Experiment-----|
Epoch [5/30], Train Loss: 1.4553, Train Accuracy: 0.4624, Test Loss: 1.4555, Test Accuracy: 0.4591
Epoch [10/30], Train Loss: 1.1787, Train Accuracy: 0.5777, Test Loss: 1.2408, Test Accuracy: 0.5511
Epoch [15/30], Train Loss: 0.8979, Train Accuracy: 0.6613, Test Loss: 1.2459, Test Accuracy: 0.5552
Epoch [20/30], Train Loss: 0.7757, Train Accuracy: 0.7284, Test Loss: 1.2323, Test Accuracy: 0.5821
Epoch [25/30], Train Loss: 0.6495, Train Accuracy: 0.7842, Test Loss: 1.2475, Test Accuracy: 0.5993
Epoch [30/30], Train Loss: 0.3189, Train Accuracy: 0.8976, Test Loss: 1.1725, Test Accuracy: 0.5933
Finished Training
Final test AccuracyL2_0.001:(0.89762667896747589, 0.59332667896747589)
```

L1 regularization with different Lambda



Comment: Model is not overfitting after L1 regularization and test and training accuracy curves are very close for every value of lambda. Best result is achieved with smaller lambda 0.001.

L2 regularization with different Lambda



Comment: for value of lambda 0.05, training and test curves are closer and hence with low variance it is not overfitting as much as compared to other values of lambda.

Exercise 3: Optimizers

In the last part, you will experiment with two different optimizers i.e. SGD and AdamOptimizer, more specifically, how robust they are to the initial learning rate. The choices for your initial learning rate are left up to you. Please compare and contrast the behavior of these two optimizers specifically on how they react when presented with different learning rates. You should plot the training curves as have been requested in the exercises above.

```
In (15): # setting device
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

in_channels = 3 # input channels of image
num_classes = 10 # number of output class [0-10]
num_epochs = 50
batch_size = 512

transform = transforms.Compose([transforms.Resize((32,32)),
                               transforms.ToTensor()])

# getting training and test data
train_dataset = datasets.CIFAR10(root='dataset/', train=True, transform=transform, download=True)
# splitting dataset into half with random indexes
indexList = random.sample(list(range(0,len(train_dataset))), (len(train_dataset)//2))
train_dataset_half = torch.utils.data.Subset(train_dataset, indexList)

test_dataset = datasets.CIFAR10(root='dataset/', train=False, transform=transform, download=True)
train_loader = DataLoader(dataset=train_dataset_half, batch_size=batch_size, shuffle=True)
test_loader = DataLoader(dataset=test_dataset, batch_size=batch_size, shuffle=True)

# different setting of learning rates
learning_rate=0.001
n_total_steps = len(train_loader)
criterion = nn.NLLLoss()
grid = (('Adam', 0.001), ('Adam', 0.0001), ('SGD', 0.5), ('SGD', 0.1), ('SGD', 0.01))

for opt, lr in grid:
    # Initialize network
    model = CNN(input_channels=in_channels, num_of_classes=num_classes).to(device)

    # Initializing Optimizer
    if (opt=='Adam'):
        optimizer= torch.optim.Adam(model.parameters(), lr=lr)
    else:
        optimizer= torch.optim.SGD(model.parameters(), lr=lr)

    # Initializing writer for tensorboard
    desc = '%s' % str(opt) + '%s' % str(lr)

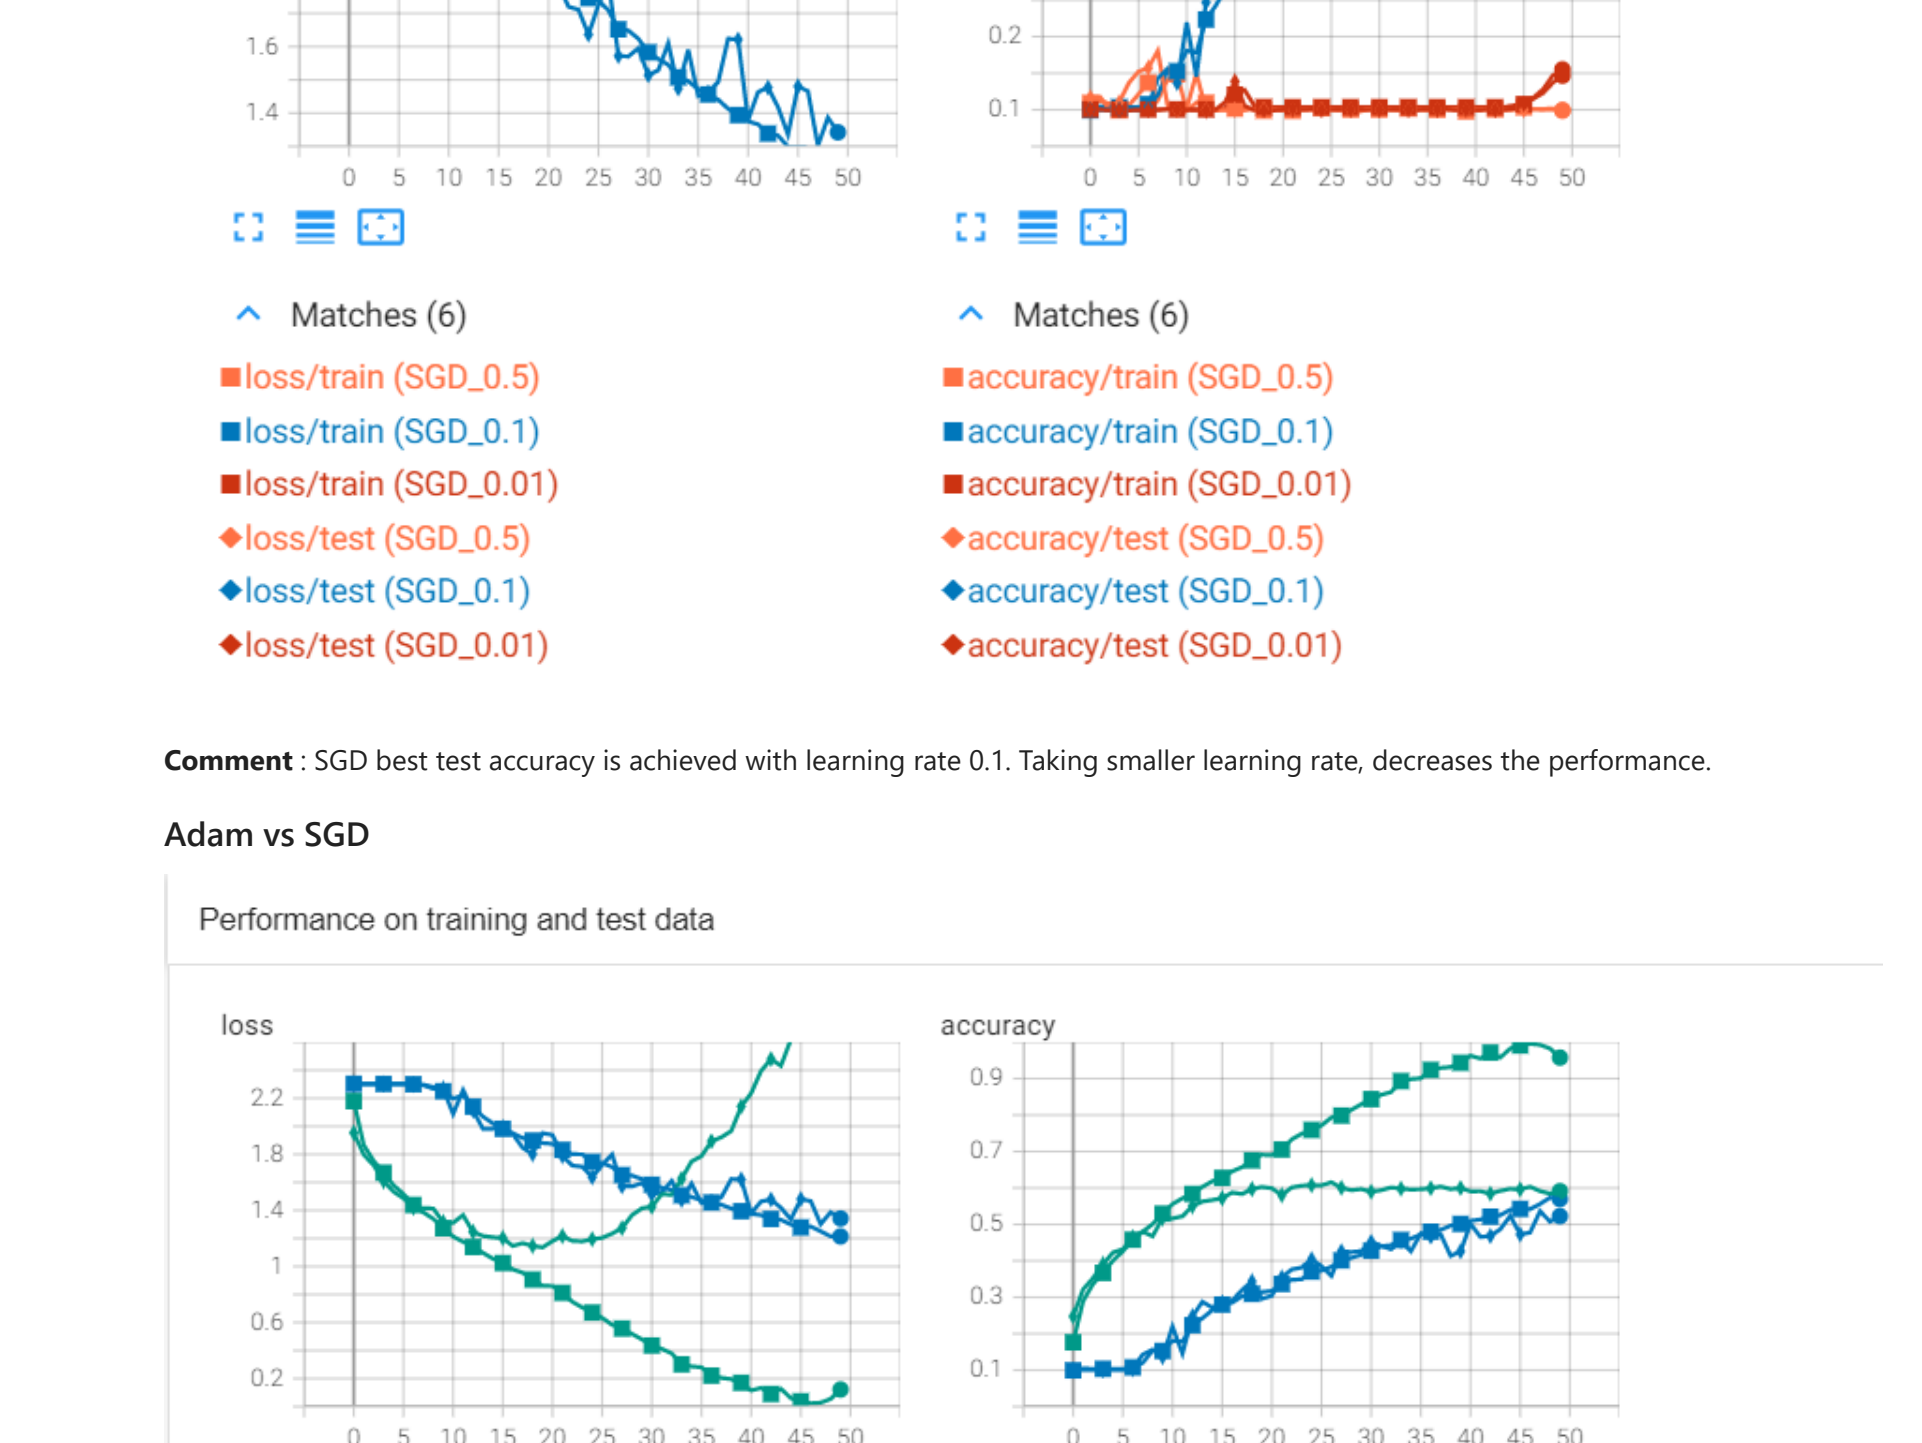
    print('----- Experiment with Optimizer = %s---Learning Rate = %s-----' % (opt, lr))
    print('-----')

    model = trainModel(model, optimizer, criterion, learning_rate, train_loader, test_loader,
    num_epochs, desc)

    writer = None
    epoch = None
    totalAccuracy = evalModel(model, optimizer, criterion, test_loader, epoch, writer)
    print('Final test Accuracy'+desc+'%s' % str(totalAccuracy))

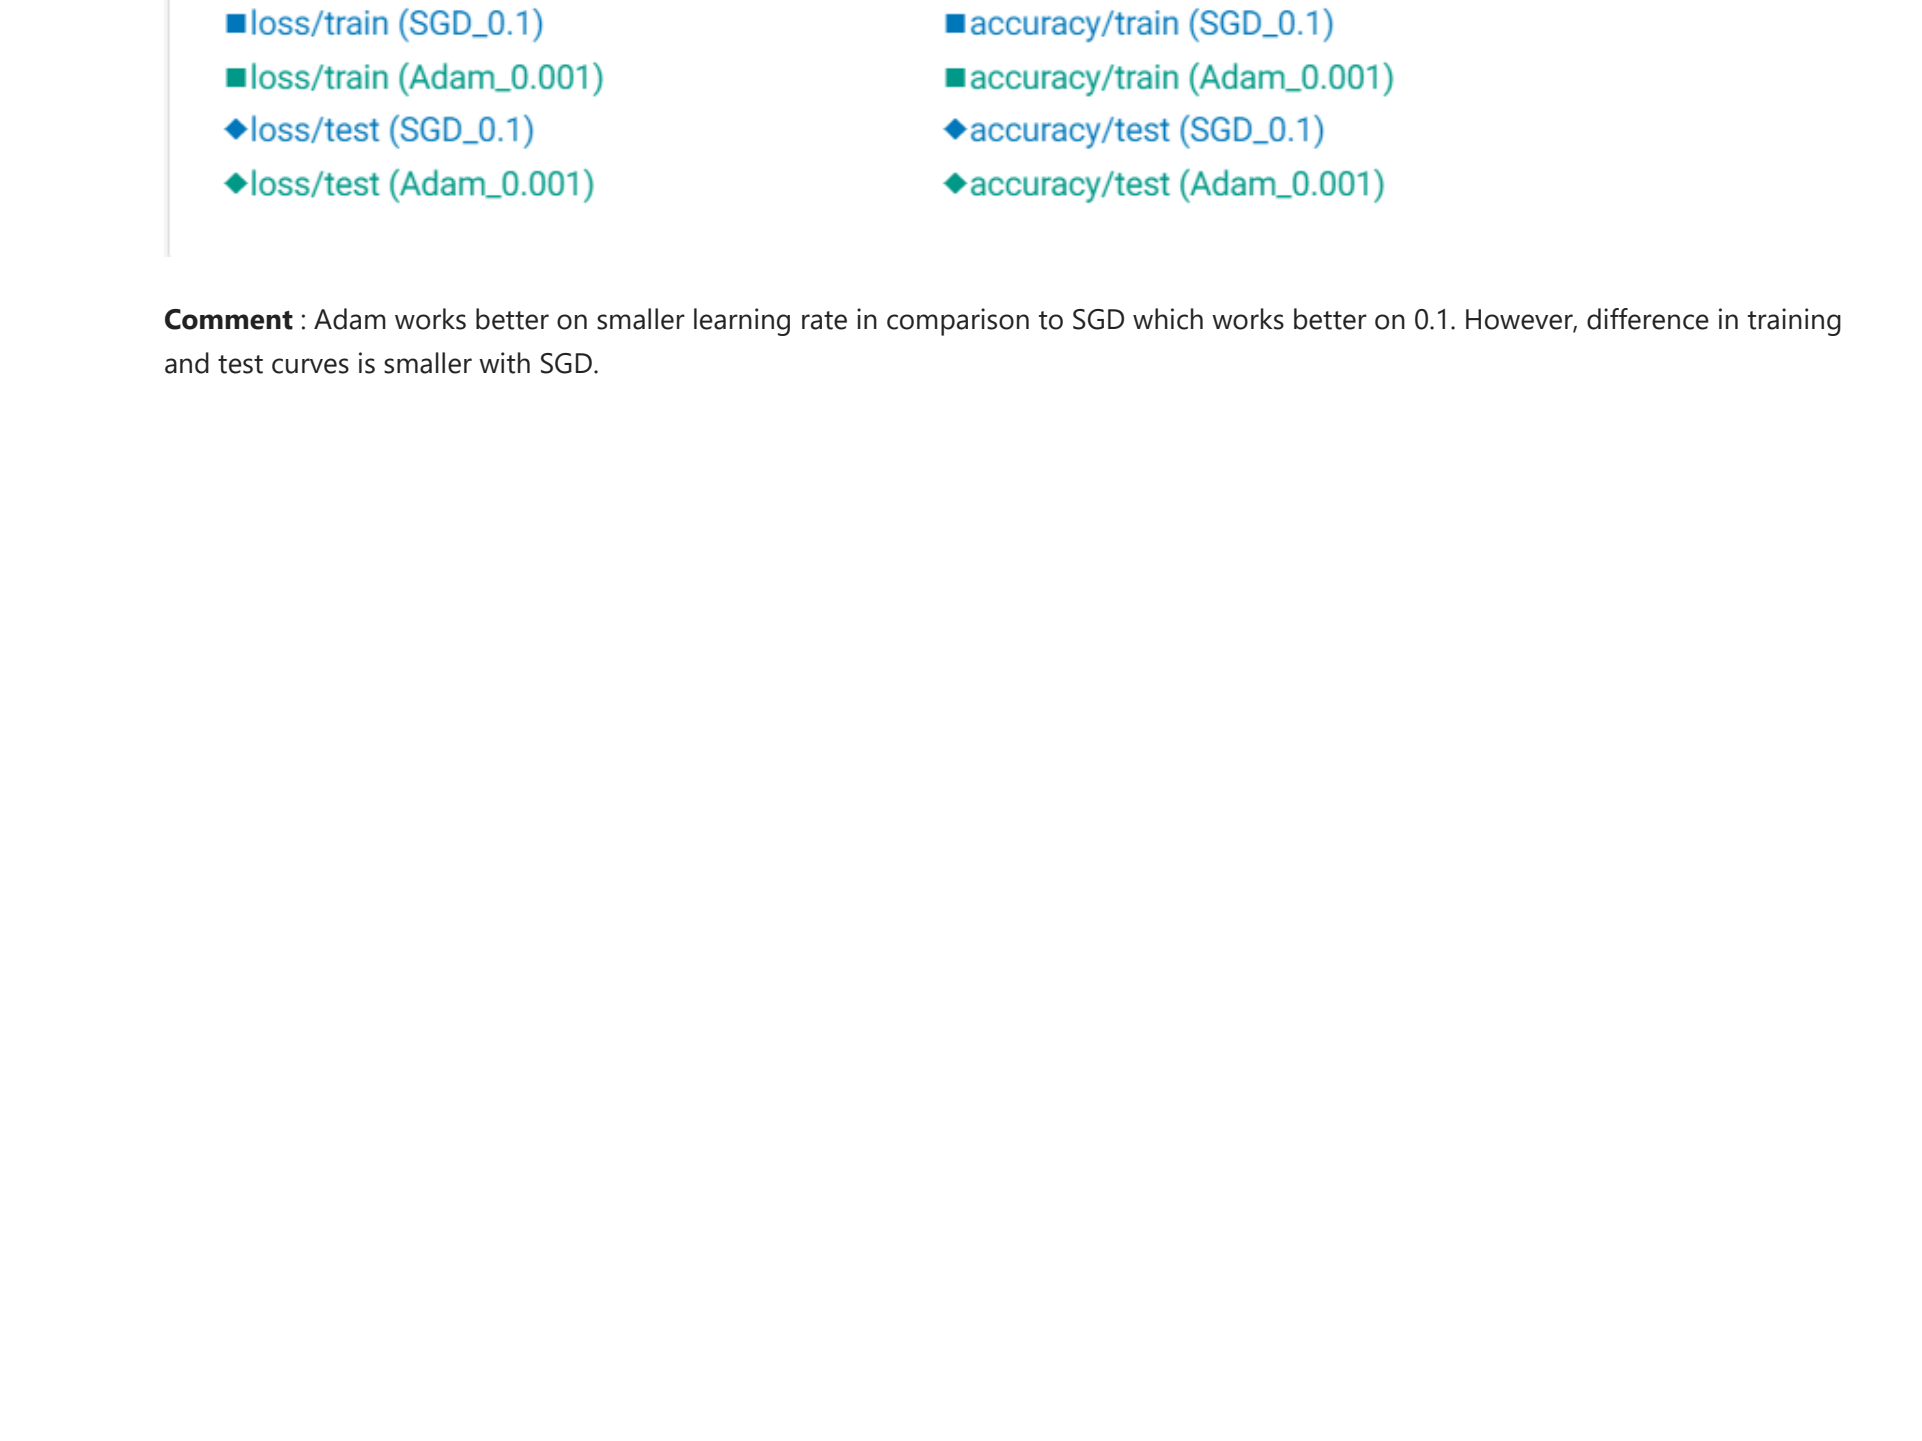
Files already downloaded and verified
Files already downloaded and verified
|----- Experiment with Optimizer = Adam---Learning Rate = 0.001-----|
Epoch [5/50], Train Loss: 1.9263, Train Accuracy: 0.2834, Test Loss: 1.8832, Test Accuracy: 0.3098
Epoch [10/50], Train Loss: 1.7809, Train Accuracy: 0.3401, Test Loss: 1.7784, Test Accuracy: 0.3481
Epoch [15/50], Train Loss: 1.6168, Train Accuracy: 0.3847, Test Loss: 1.7311, Test Accuracy: 0.3567
Epoch [20/50], Train Loss: 1.5708, Train Accuracy: 0.4057, Test Loss: 1.7143, Test Accuracy: 0.3576
Epoch [25/50], Train Loss: 1.5212, Train Accuracy: 0.4456, Test Loss: 1.6899, Test Accuracy: 0.3942
Epoch [30/50], Train Loss: 1.4747, Train Accuracy: 0.4807, Test Loss: 1.6707, Test Accuracy: 0.4081
Epoch [35/50], Train Loss: 1.3828, Train Accuracy: 0.4972, Test Loss: 1.6027, Test Accuracy: 0.3951
Epoch [40/50], Train Loss: 1.2907, Train Accuracy: 0.5281, Test Loss: 1.5013, Test Accuracy: 0.3819
Epoch [45/50], Train Loss: 1.2893, Train Accuracy: 0.5537, Test Loss: 1.3195, Test Accuracy: 0.3908
Epoch [50/50], Train Loss: 1.1645, Train Accuracy: 0.5763, Test Loss: 2.0987, Test Accuracy: 0.3824
Finished Training
Final test AccuracyAdam_0.001:(0.3822667896747589, 0.29332537651062)
|----- Experiment with Optimizer = Adam---Learning Rate = 0.0001-----|
Epoch [5/50], Train Loss: 1.8451, Train Accuracy: 0.2683, Test Loss: 1.8195, Test Accuracy: 0.3177
Epoch [10/50], Train Loss: 1.6855, Train Accuracy: 0.3182, Test Loss: 1.7078, Test Accuracy: 0.3668
Epoch [15/50], Train Loss: 1.5531, Train Accuracy: 0.3614, Test Loss: 1.6139, Test Accuracy: 0.3567
Epoch [20/50], Train Loss: 1.3570, Train Accuracy: 0.4057, Test Loss: 1.5143, Test Accuracy: 0.3576
Epoch [25/50], Train Loss: 1.2139, Train Accuracy: 0.4456, Test Loss: 1.4143, Test Accuracy: 0.3942
Epoch [30/50], Train Loss: 1.0708, Train Accuracy: 0.4807, Test Loss: 1.3195, Test Accuracy: 0.4081
Epoch [35/50], Train Loss: 0.9309, Train Accuracy: 0.5281, Test Loss: 1.2196, Test Accuracy: 0.3867
Epoch [40/50], Train Loss: 0.7999, Train Accuracy: 0.5656, Test Loss: 1.1470, Test Accuracy: 0.4993
Epoch [45/50], Train Loss: 0.6495, Train Accuracy: 0.5867, Test Loss: 1.2475, Test Accuracy: 0.5993
Epoch [50/50], Train Loss: 0.3189, Train Accuracy: 0.8976, Test Loss: 1.1725, Test Accuracy: 0.5933
Finished Training
Final test AccuracyAdam_0.0001:(0.59332667896747589, 0.29332537651062)
|----- Experiment with Optimizer = SGD---Learning Rate = 0.5-----|
Epoch [5/50], Train Loss: 1.8048, Train Accuracy: 0.3242, Test Loss: 1.7867, Test Accuracy: 0.3346
Epoch [10/50], Train Loss: 1.6535, Train Accuracy: 0.3839, Test Loss: 1.6549, Test Accuracy: 0.3760
Epoch [15/50], Train Loss: 1.4859, Train Accuracy: 0.4231, Test Loss: 1.5143, Test Accuracy: 0.4120
Epoch [20/50], Train Loss: 1.2489, Train Accuracy: 0.4503, Test Loss: 1.4833, Test Accuracy: 0.4490
Epoch [25/50], Train Loss: 1.1374, Train Accuracy: 0.4717, Test Loss: 1.4135, Test Accuracy: 0.4261
Epoch [30/50], Train Loss: 1.0708, Train Accuracy: 0.4807, Test Loss: 1.3195, Test Accuracy: 0.4081
Epoch [35/50], Train Loss: 0.9309, Train Accuracy: 0.5281, Test Loss: 1.2196, Test Accuracy: 0.3867
Epoch [40/50], Train Loss: 0.7999, Train Accuracy: 0.5656, Test Loss: 1.1470, Test Accuracy: 0.4993
Epoch [45/50], Train Loss: 0.6495, Train Accuracy: 0.5867, Test Loss: 1.2475, Test Accuracy: 0.5993
Epoch [50/50], Train Loss: 0.3189, Train Accuracy: 0.8976, Test Loss: 1.1725, Test Accuracy: 0.5933
Finished Training
Final test AccuracySGD_0.5:(0.59332667896747589, 0.29332537651062)
|----- Experiment with Optimizer = SGD---Learning Rate = 0.1-----|
Epoch [5/50], Train Loss: 1.8048, Train Accuracy: 0.3242, Test Loss: 1.7867, Test Accuracy: 0.3346
Epoch [10/50], Train Loss: 1.6535, Train Accuracy: 0.3839, Test Loss: 1.6549, Test Accuracy: 0.3760
Epoch [15/50], Train Loss: 1.4859, Train Accuracy: 0.4231, Test Loss: 1.5143, Test Accuracy: 0.4120
Epoch [20/50], Train Loss: 1.2489, Train Accuracy: 0.4503, Test Loss: 1.4833, Test Accuracy: 0.4490
Epoch [25/50], Train Loss: 1.1374, Train Accuracy: 0.4717, Test Loss: 1.4135, Test Accuracy: 0.4261
Epoch [30/50], Train Loss: 1.0708, Train Accuracy: 0.4807, Test Loss: 1.3195, Test Accuracy: 0.4081
Epoch [35/50], Train Loss: 0.9309, Train Accuracy: 0.5281, Test Loss: 1.2196, Test Accuracy: 0.3867
Epoch [40/50], Train Loss: 0.7999, Train Accuracy: 0.5656, Test Loss: 1.1470, Test Accuracy: 0.4993
Epoch [45/50], Train Loss: 0.6495, Train Accuracy: 0.5867, Test Loss: 1.2475, Test Accuracy: 0.5993
Epoch [50/50], Train Loss: 0.3189, Train Accuracy: 0.8976, Test Loss: 1.1725, Test Accuracy: 0.5933
Finished Training
Final test AccuracySGD_0.1:(0.59332667896747589, 0.29332537651062)
|----- Experiment with Optimizer = SGD---Learning Rate = 0.01-----|
Epoch [5/50], Train Loss: 1.8048, Train Accuracy: 0.3242, Test Loss: 1.7867, Test Accuracy: 0.3346
Epoch [10/50], Train Loss: 1.6535, Train Accuracy: 0.3839, Test Loss: 1.6549, Test Accuracy: 0.3760
Epoch [15/50], Train Loss: 1.4859, Train Accuracy: 0.4231, Test Loss: 1.5143, Test Accuracy: 0.4120
Epoch [20/50], Train Loss: 1.2489, Train Accuracy: 0.4503, Test Loss: 1.4833, Test Accuracy: 0.4490
Epoch [25/50], Train Loss: 1.1374, Train Accuracy: 0.4717, Test Loss: 1.4135, Test Accuracy: 0.4261
Epoch [30/50], Train Loss: 1.0708, Train Accuracy: 0.4807, Test Loss: 1.3195, Test Accuracy: 0.4081
Epoch [35/50], Train Loss: 0.9309, Train Accuracy: 0.5281, Test Loss: 1.2196, Test Accuracy: 0.3867
Epoch [40/50], Train Loss: 0.7999, Train Accuracy: 0.5656, Test Loss: 1.1470, Test Accuracy: 0.4993
Epoch [45/50], Train Loss: 0.6495, Train Accuracy: 0.5867, Test Loss: 1.2475, Test Accuracy: 0.5993
Epoch [50/50], Train Loss: 0.3189, Train Accuracy: 0.8976, Test Loss: 1.1725, Test Accuracy: 0.5933
Finished Training
Final test AccuracySGD_0.01:(0.59332667896747589, 0.29332537651062)
```

Adam with different learning rates




Comment: Adam's test accuracy is achieved with learning rate 0.001 however difference b/w train and test accuracy is also high for learning rate 0.001. With smaller learning rate 0.0001 this difference is small. However Adam is quite robust to learning rate.

SGD with different learning rates



Comment: SGD best test accuracy is achieved with learning rate 0.1. Taking smaller learning rate, decreases the performance.

Adam vs SGD



Comment: Adam works better on smaller learning rate in comparison to SGD which works better on 0.1. However, difference in training and test curves is smaller with SGD.